
DuoGPT: Training-free Dual Sparsity through Activation-aware Pruning in LLMs

Ruokai Yin, Yuhang Li, Donghyun Lee, Priyadarshini Panda
Yale University
ruokai.yin@yale.edu

Abstract

Large language models (LLMs) deliver strong performance but are difficult to deploy due to high memory and compute costs. While pruning reduces these demands, most methods ignore activation sparsity observed at runtime. We reinterpret activation sparsity as dynamic structured weight sparsity and propose **DuoGPT**, a unified framework that constructs dual-sparse (spMspV) workloads by combining unstructured weight pruning with activation sparsity. To preserve accuracy, we extend the Optimal Brain Compression (OBC) framework with activation-aware calibration and introduce output residuals from the dense model as correction terms. We further optimize the solution for efficient GPU execution, enabling scalability to billion-parameter LLMs. Evaluations on LLaMA-2 and LLaMA-3 show that **DuoGPT** outperforms state-of-the-art structured pruning methods by up to 9.17% accuracy at an iso-speedup of $1.39\times$ compared to the baseline dense model. Code is available at [Github](#).

1 Introduction

Large language models (LLMs) have made significant advances in performance across a range of complex, real-world language tasks (Brown et al. 2020, Touvron et al. 2023). However, their massive parameter counts present practical challenges for deploying these large pre-trained models during inference. For example, deploying LLaMA-2-70B (Touvron et al. 2023), an open-source pre-trained LLM with 70 billion parameters, typically requires 150 GB of GPU and CPU RAM, along with 60 GFLOPs of computation to decode a single token. Consequently, there is widespread interest in reducing both the storage and computational requirements of LLMs through techniques collectively referred to as model compression. Existing model compression techniques for LLMs include quantization (Frantar et al. 2022, Ashkboos et al. 2024b), tensor decomposition (Wang et al. 2024, Lin et al. 2024), and pruning (Frantar & Alistarh 2023, Ashkboos et al. 2024a). This work focuses on pruning, specifically incorporating activation sparsity into the one-shot weight pruning framework.

One-shot pruning methods selectively zero out the elements in the weight matrices of an LLM during one forward pass, and optionally update the remaining non-zero weights to compensate for the pruning error. Depending on the freedom in the zero patterns, the pruning can be further categorized into *unstructured pruning* and *structured pruning*. In general, the following tradeoff between the two categories long stands: more structured sparsity yields more speedup¹, while more unstructured sparsity associates with better accuracy performance. In this work, we strike a balance between the two by incorporating activation sparsity. We observe that during the decoding stage of LLMs (i.e., General Matrix-Vector (GEMV) operations²), activation sparsity can be interpreted as structured weight sparsity, as only weight rows corresponding to non-zero activations are involved

¹We focus on the speedup of running LLMs on general purpose architectures, e.g., GPUs.

²We primarily focus on single-batch decoding in this work.




Figure 1: GEMV operation for the single-batch decoding stage under different types of sparsity.

in computation. Further, activation sparsity has been observed in most LLMs due to the non-linear activation functions (Liu et al. 2023, Li et al. 2022). Despite its acceleration potential, activation sparsity still requires storing the full dense model on the GPU, resulting in significant memory overhead. To reduce this burden, we propose leveraging activation sparsity on top of unstructured weight pruning. Our method thus produces dual-sparse LLM workloads that achieve a favorable balance between accuracy and efficiency, as illustrated in Figure 1.

Although combining dynamic activation sparsity with static weight sparsity is appealing, doing so naively can degrade performance. To address this, we extend the Optimal Brain Compression (OBC) framework (Frantar & Alistarh 2023) to compensate for runtime errors introduced by activation sparsity. Inspired by asymmetric calibration (Li et al. 2025), we reconstruct a layerwise objective using the sparse activation from the pruned model and the corresponding dense activation from the original model. This enables weight updates during pruning calibration, mitigating accuracy degradation caused by sparse activations. Our contributions are as follows:

1. We reinterpret activation sparsity as a dynamic, structured form of weight sparsity. Combined with unstructured pruning, this perspective enables the construction of sparse matrix–sparse vector (spMspV) workloads for compressing and accelerating LLMs during decoding.
2. To compensate for the runtime errors introduced by dynamic activation sparsity, we propose **DuoGPT**, a layerwise iterative pruning method that extends the OBC framework to activation-sparsity-aware unstructured pruning in LLMs. With an efficient implementation, **DuoGPT** can calibrate a 70B-parameter LLaMA-3 model in under 130 minutes on a single A100 80GB GPU.
3. We conduct extensive experiments on LLaMA-2 and 3, demonstrating that **DuoGPT** significantly improves dual-sparse model performance across tasks. Under 50% dual weight-activation sparsity, **DuoGPT** outperforms state-of-the-art structured-pruning (e.g., ShortGPT (Men et al. 2024)) by 9.17% accuracy at an iso-speedup of roughly $1.4\times$ relative to the dense baseline. Compared to the sparse-activation (e.g., R-Sparse (Zhang et al. 2025)) methods, **DuoGPT** achieves up to a $1.97\times$ model size reduction at iso-accuracy.

2 Related Work

Activation Sparsity. Activation sparsity is a well-established property in ReLU-based LLMs (Li et al. 2022, Mirzadeh et al. 2023) and also emerges with newer activations like SwiGLU (Zhang et al. 2024a, Liu et al. 2023, Zhang et al. 2021, Haziza et al. 2025), naturally inducing sparsity in Multi-layer Perceptron (MLP) activations. This has motivated methods such as MoE-style MLP reconfiguration (Zhang et al. 2021, 2024b, Liu et al. 2023) and threshold-based activation pruning from calibration data (Lee et al. 2024). Recent efforts have extended activation pruning to all transformer layers using thresholding (Liu et al. 2025) or hybrid SVD-based techniques (Zhang et al. 2025). Unlike these approaches—which retain dense weights—our work integrates activation sparsity with unstructured weight pruning, yielding a dual-sparse LLM workload. Notably, threshold-based techniques from prior work are complementary and can be incorporated into our framework.

One-shot Weight Pruning for LLMs. One-shot pruning, originally proposed by Han et al. (2015), removes weights based on magnitude but performs poorly at LLM scale (Frantar & Alistarh 2023). SparseGPT (Frantar & Alistarh 2023) improves this with an approximate regression solver, enabling accurate pruning of 100B+ parameter models on a single GPU without fine-tuning. SlimGPT (Ling

et al. 2024) extends this to structured pruning. Wanda (Sun et al. 2023) offers compensation-free unstructured pruning using magnitude of weight and activation. SliceGPT (Ashkboos et al. 2024a) prunes both weights and activations while preserving computational invariance, but suffers sharp accuracy drops beyond 30% sparsity. Other approaches prune at the block or submodule level using importance scores (Men et al. 2024, Zhong et al. 2024, Sandri et al. 2025).

Optimal Brain Compression (OBC). OBC builds on the Optimal Brain Surgeon (OBS) framework (LeCun et al. 1989, Hassibi & Stork 1992), which leverages second-order derivatives and the inverse Hessian to guide weight pruning. OBC (Frantar & Alistarh 2022) approximates this approach using an iterative solver, yielding closed-form solutions for optimal weight selection and updates. SparseGPT (Frantar & Alistarh 2023) was the first to scale this framework effectively to LLMs. Our work extends this line by incorporating activation sparsity into the OBC formulation and introducing a correction term derived from dense inputs. This design is inspired by the asymmetric calibration strategy recently proposed by Li et al. (2025) for OBC-based quantization.

3 Preliminaries

Notations. Throughout this paper, row vectors are denoted by bold lowercase letters and matrices by bold uppercase letters. For example, $\mathbf{w} \in \mathcal{R}^{1 \times k}$ represents a row vector of the weight. Consequently, the linear operation between the weight $\mathbf{W} \in \mathcal{R}^{n \times k}$ and the input calibration matrices $\mathbf{X} \in \mathcal{R}^{k \times m}$ is expressed as: $\mathbf{Y} = \mathbf{W}\mathbf{X}$ ³. Here, k denotes the hidden dimension, n is the number of output channels, and m is the number of input tokens. Pruning is done through the element-wise multiplication, for example, $\hat{\mathbf{w}} = \mathbf{w} \odot \mathbf{m}^{\mathbf{w}}$ represents the pruning of one row of weight. Here, $\mathbf{m} \in \{0, 1\}$ is the bitmask. We use superscripts to specify different masks. Subscripts are used to specify the subsets of vectors and matrices. A negative subscript index on a matrix indicates the removal of the corresponding row and column. For example, \mathbf{A}_{-j} is the matrix \mathbf{A} with its j^{th} column and j^{th} row removed.

Activation Sparsity. We induce activation sparsity by element-wise multiplying the input activation \mathbf{x} with a binary mask $\mathbf{m}^{\mathbf{x}}$, i.e., $\mathbf{x} \odot \mathbf{m}^{\mathbf{x}}$. The activation sparsity level $p^{\mathbf{x}}$ is defined as the proportion of zeros in $\mathbf{m}^{\mathbf{x}}$. For a given linear layer l during the decoding stage, the computation becomes: $\mathbf{y} = \mathbf{W}(\mathbf{x} \odot \mathbf{m}^{\mathbf{x}}) = \sum_{i \in \mathbf{m}_i^{\mathbf{x}}=1} \mathbf{W}_{:,i} \hat{\mathbf{x}}_i$, where $\hat{\mathbf{x}}_i$ denotes the sparse input vector. During calibration, when the input calibration data is in matrix form, we enforce that each column of the input has a sparsity level of $p^{\mathbf{x}}$. Activations are pruned based on their absolute magnitude on the fly. In contrast to prior works (Liu et al. 2025, Zhang et al. 2025) that search for layer-specific optimal activation sparsity levels, we investigate a simpler setting in which uniform sparsity is applied across all transformer layers. At decoding time, our method is agnostic to how activation sparsity is induced (Details can be found in Section 5).

OBC Framework for Pruning. The OBC (Frantar & Alistarh 2022) pruning framework converts dense weight matrices \mathbf{W} to a sparse one $\hat{\mathbf{W}}$ through a calibration process to preserve model performance. This calibration process is efficiently implemented by SparseGPT (Frantar & Alistarh 2023) as an iterative layer-wise pruning framework for LLMs. Formally, for each row of weight matrix, the calibration targets to minimize the differences between the original and pruned layer outputs: $\min_{\Delta \mathbf{w}} \|(\Delta \mathbf{w} + \mathbf{w})\mathbf{X} - \mathbf{w}\mathbf{X}\|_F^2$, s.t. $\Delta \mathbf{w} \mathbf{e}_p^{\top} + \mathbf{w}_p = 0$, where \mathbf{e}_p^{\top} is the unit vector that corresponds to \mathbf{w}_p , the weight that got removed. The loss for \mathbf{w}_p and the updates to remaining weights are:

$$\mathcal{L}_p = \frac{\mathbf{w}_p^2}{\mathbf{H}_{pp}^{-1}}, \quad \Delta \mathbf{w} = -\frac{\mathbf{w}_p}{\mathbf{H}_{pp}^{-1}} \cdot \mathbf{H}_{p,:}^{-1}. \quad (1)$$

Here $\mathbf{H}^{-1} = (\mathbf{X}\mathbf{X}^{\top})^{-1}$ denotes the inverse Hessian matrix. The weights are sorted by their loss value and the weight with lowest loss value is removed. Then, we compensate the remaining weights according to Equation 1. Once the p^{th} weight is removed, the inverse Hessian is updated through Gaussian elimination: $\mathbf{H}_{-p}^{-1} = (\mathbf{H}^{-1} - \frac{\mathbf{H}_{:,p}^{-1}\mathbf{H}_{p,:}^{-1}}{\mathbf{H}_{pp}^{-1}})$. This process iterates until the preset $p^{\mathbf{w}}$ layerwise weight sparsity is achieved. The resulted sparse model follows the closed-form solution.

³We align the same notation with prior work in our method description. In implementation or illustration, the weight matrices are stored in column-major format, as shown in Figure. 1 and 2(a).

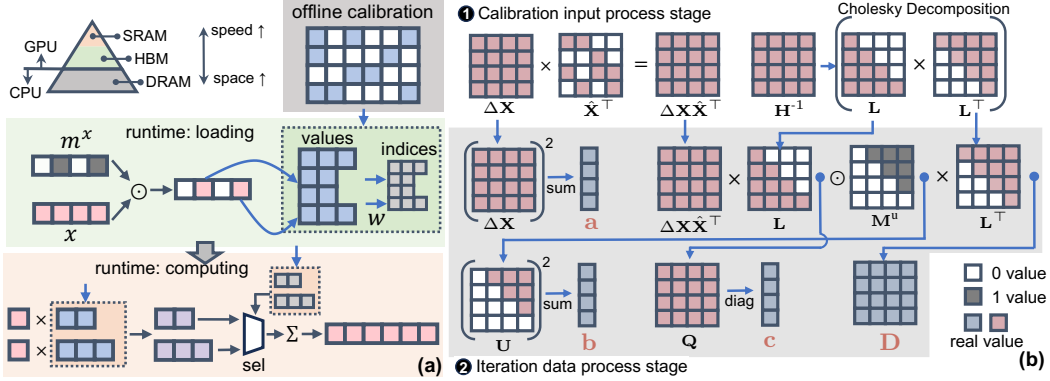


Figure 2: (a) Illustration of how dual-sparsity accelerates the decoding stage of LLMs by saving computation, memory loading, and storage. (b) Computing paradigm of the **DuoGPT**'s efficient GPU implementation. We neglect the element-wise division of $\text{diag}(\mathbf{L})$ for \mathbf{c} in the figure.

4 DuoGPT

4.1 Activation Sparsity in Sparse LLMs

By leveraging activation sparsity, only certain rows of weights are fetched from High Bandwidth Memory (HBM) into Static RAM (SRAM) and further into the computation cores. However, since it is not known beforehand which specific rows will be required at runtime, memory storage cannot be reduced. As a result, the entire dense model must still be fully materialized in the GPU's HBM.

To address this limitation, we propose introducing sparsity into both activations and weights—a strategy referred to as **dual-sparsity**. Specifically, we prune the dense model into a sparse version during an offline calibration process. During inference, only the compressed model parameters are loaded and transferred from CPU to GPU, significantly reducing both GPU HBM bandwidth and memory storage requirements. During linear layer computations, only the rows of weight matrices corresponding to non-zero activations are fetched from HBM into SRAM. Moreover, because each weight row is stored in compressed format, additional savings are achieved in SRAM loading, SRAM storage, and GEMM core computation. The overall runtime execution flow for this dual-sparse GEMV workload is illustrated in Figure 2(a).

A potential challenge with dual-sparsity lies in the overlap between activation and weight sparsity. Given activation and weight sparsity levels of p^x and p^w , respectively, the fraction of weights loaded from HBM to SRAM could approach $(1 - p^x)$ in the worst case—particularly if the weight sparsity distribution is highly skewed. Empirically, however, we found that incorporating activation sparsity into the pruning calibration process helps distribute weight sparsity more uniformly across rows. For instance, for our calibrated LLaMA-2-7B model with $p^x = p^w = 0.5$, the maximum fraction of weights loaded from HBM to SRAM is approximately 0.25 (More details in Section 5 and Appendix). The remaining challenge is to determine *how activation sparsity can be effectively integrated into the weight pruning calibration process*.

4.2 Activation Sparsity aware Pruning Calibration

A central challenge in leveraging activation sparsity for model pruning lies in incorporating this sparsity pattern directly into the calibration process. This section introduces a principled framework for solving this problem. Let \mathbf{X} denote the output from the previously pruned layer during calibration. Applying magnitude-based pruning to \mathbf{X} yields a sparse version of the calibration input, denoted as $\tilde{\mathbf{X}}$, which maintains uniform sparsity p^x across its columns. The calibration objective thus becomes minimizing the expression $\|(\Delta \mathbf{w} + \mathbf{w})\tilde{\mathbf{X}} - \mathbf{w}\hat{\mathbf{X}}\|_F^2$. While utilizing sparse calibration data helps the model quickly adapt to activation sparsity during inference, it inevitably reduces the information density available for accurately compensating the remaining unpruned weights. This fundamental trade-off necessitates a more sophisticated approach.

The solution lies in asymmetric calibration (Li et al. 2025), which incorporates both sparse and dense data to mitigate information loss. Given $\tilde{\mathbf{X}}$ as the dense output from the dense model, the target

becomes $\|(\Delta \mathbf{w} + \mathbf{w})\hat{\mathbf{X}} - \mathbf{w}\tilde{\mathbf{X}}\|_F^2$. Setting $\mathbf{w}(\tilde{\mathbf{X}} - \hat{\mathbf{X}}) = \mathbf{r}$ as the output residual between dense and sparse data, the target simplifies to $\|\Delta \mathbf{w}\hat{\mathbf{X}} - \mathbf{r}\|_F^2$. With the constraint of $\Delta \mathbf{w}\mathbf{e}_p^\top + \mathbf{w} = 0$, the Lagrangian formulation becomes:

$$\min_{\lambda, \Delta \mathbf{w}} \mathcal{L}_p = \|\Delta \mathbf{w}\hat{\mathbf{X}} - \mathbf{r}\|_F^2 + \lambda(\Delta \mathbf{w}\mathbf{e}_p^\top + \mathbf{w}_p). \quad (2)$$

Solving this Lagrangian with $\mathbf{H} = \hat{\mathbf{X}}\hat{\mathbf{X}}^\top$ yields the induced error and optimal weight updates for the p -th weight (if removed):

$$\mathcal{L}_p = \frac{\mathbf{w}_p^2}{\mathbf{H}_{pp}^{-1}} + \mathbf{r}\mathbf{r}^\top - \mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}}\mathbf{r}^\top + \frac{2\mathbf{w}_p}{\mathbf{H}_{pp}^{-1}} \mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}_{:,p}^{-1}, \quad \Delta \mathbf{w} = -\frac{\mathbf{w}_p}{\mathbf{H}_{pp}^{-1}} \cdot \mathbf{H}_{p,:}^{-1} + \mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1}. \quad (3)$$

Equation 3 defines the iterative framework, **DuoGPT**, which prunes and calibrates the model using a closed-form solution. For each row \mathbf{w} , the algorithm identifies the p -th weight for pruning by computing $p = \arg \min_p (\mathcal{L}_p)$, then calculates a compensation term $\Delta \mathbf{w}$ to update the remaining weights. This process repeats until the target pruning ratio p^* is achieved.

Despite its theoretical soundness of effectively incorporating activation sparsity ($\hat{\mathbf{X}}$) into the calibration process and compensating the associated information loss with dense model’s information ($\tilde{\mathbf{X}}$), the direct implementation of this framework becomes computationally infeasible for large-scale LLMs with billions of parameters. Two critical bottlenecks emerge: (1) A separate Hessian inverse (\mathbf{H}^{-1}) must be maintained and updated for each weight row at every iteration (Frantar & Alistarh 2023). (2) Selecting which weight to prune requires computing the loss \mathcal{L} in Equation 3 for every candidate weight, followed by recalculating the corresponding compensation term $\Delta \mathbf{w}$ after pruning.

The naive implementation results in a total computational complexity of $\mathcal{O}(nmk^2 + nk^3)$ for a single layer. In the context of LLM calibration, both m and k reach substantial values (e.g., $m = 128 \times 2048 = 262,144$ and $k = 4096$ for LLaMA-2-7B), far exceeding the practical limits of modern GPUs. An efficient implementation is therefore essential, as detailed in the following section.

4.3 Efficient DuoGPT Algorithm

To overcome these challenges, a simplification is first made: assume that all weights to be pruned have already been selected, i.e., a fixed pruning mask \mathbf{M}^w is provided. Under this assumption, the asynchronous Hessian issue can be addressed through Hessian synchronization and weight freezing, as introduced in SparseGPT (Frantar & Alistarh 2023). The key insight is to process all rows in parallel by iterating through columns in a fixed order, rather than strictly following the optimal pruning order. This synchronization enables all rows to share a common Hessian, requiring only a single copy of \mathbf{H}^{-1} in GPU memory. Moreover, all Hessian updates can be precomputed in a single step using Cholesky decomposition, where $\mathbf{H}^{-1} = \mathbf{L}\mathbf{L}^\top$ and \mathbf{L} is the lower-triangular Cholesky factor. This formulation resolves the asynchronous Hessian challenge with minimal computational overhead while enhancing numerical stability (Frantar & Alistarh 2023).

The critical remaining problem is: **how to determine the pruning mask \mathbf{M}^w ?** More precisely, a pruning score metric \mathbf{S} is needed for all weights to derive the pruning mask, effectively decoupling mask selection from the pruning and calibration stages. This score must satisfy three essential criteria: (1) **Efficiency**: The score must be inexpensive to compute, expressible in vectorized form, and amenable to precomputation. (2) **Adaptivity**: It must incorporate update information across iterations to avoid the performance degradation associated with static pruning structures. (3) **Fidelity**: The score should closely approximate the original loss metric \mathcal{L} to ensure that the final result remains a good approximation of the exact closed-form solution.

To begin, the pruning score \mathbf{S} is constructed using the error \mathcal{L} incurred when each weight is removed, as derived from Equation 3. For each weight column p , the pruning scores are defined as:

$$\mathbf{S}_{:,p} = \frac{\mathbf{W}_{:,p}^2}{\mathbf{H}_{pp}^{-1}} + \text{diag}(\mathbf{R}\mathbf{R}^\top) - \text{diag}(\mathbf{R}\hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}}\mathbf{R}^\top) + \frac{2\mathbf{W}_{:,p}}{\mathbf{H}_{pp}^{-1}} \mathbf{R}\hat{\mathbf{X}}^\top \mathbf{H}_{:,p}^{-1}, \quad (4)$$

where $\mathbf{R} = \mathbf{W}(\hat{\mathbf{X}} - \tilde{\mathbf{X}})$ denotes the output residual for the full weight matrix. An observation is that Equation 4 comprises orthogonal error terms capturing distinct aspects of the objective function’s landscape. The first term, inherited from SparseGPT, quantifies the intrinsic information loss from

removing each weight in isolation. The subsequent terms emerge from **DuoGPT** and operate in complementary subspaces. The second term captures the quadratic self-interaction of output residuals, representing errors from the sparse-dense activation discrepancy. The third term functions as a negative correction, quantifying how each weight’s removal affects activation-sparsity-induced errors. The final term represents a cross-correlation between weight importance and activation-sparsity effects, capturing dependencies that enable adaptive compensation for dual sparsity effects. Together, these terms form a holistic metric that precisely quantifies each weight’s contribution to model performance in the dual-sparse regime.

An important observation emerges from Equation 4: although output residuals help compensate for information loss introduced by activation sparsity, they must be updated whenever weights are modified: $\mathbf{R} \leftarrow \mathbf{R} - \Delta\mathbf{W}(\tilde{\mathbf{X}} - \hat{\mathbf{X}})$. This requirement not only eliminates the possibility of precomputing pruning scores but introduces a computational cost of $\mathcal{O}(mnk)$ per column, rendering the naive approach prohibitively expensive. An opportunity of simplification comes from decomposing \mathbf{R} into a sum of partial-sum matrices formed by the outer product between each weight column and the corresponding row of the input difference, $\Delta\mathbf{X} = \tilde{\mathbf{X}} - \hat{\mathbf{X}}$. Specifically, the residual matrix can be approximated as $\mathbf{R} = \sum_{j=0}^k \mathbf{W}_{:,j} \Delta\mathbf{X}_{j,:}$. Since \mathbf{R} quantifies the discrepancy between sparse and dense model outputs, it can be decomposed into contributions from individual weight-activation pairs. This decomposition transforms Equation 4 into:

$$\mathbf{S}_{:,p} = \frac{\mathbf{W}_{:,p}^2}{\mathbf{H}_{pp}^{-1}} + \mathbf{R}_p \mathbf{R}_p^\top - \mathbf{W}_{:,p} \Delta\mathbf{X}_{p,:} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} \Delta\mathbf{X}_{p,:}^\top \mathbf{W}_{:,p}^\top + \frac{2\mathbf{W}_{:,p}}{\mathbf{H}_{pp}^{-1}} \mathbf{W}_{:,p} \Delta\mathbf{X}_{p,:} \hat{\mathbf{X}}^\top \mathbf{H}_{:,p}^{-1}, \quad (5)$$

where $\mathbf{R}_p = \mathbf{W}_{:,p} \Delta\mathbf{X}_{p,:}$. Furthermore, since the pruning process is independent across rows, the term $\mathbf{W}_{:,p}^2$ can be factored out from the last three terms in Equation 5. This results in the following expression for the pruning score:

$$\mathbf{S}_{:,p} = \mathbf{W}_{:,p}^2 \left(\frac{1}{\mathbf{H}_{pp}^{-1}} + \Delta\mathbf{X}_{p,:} \Delta\mathbf{X}_{p,:}^\top - \Delta\mathbf{X}_{p,:} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} \Delta\mathbf{X}_{p,:}^\top + \frac{2}{\mathbf{H}_{pp}^{-1}} \Delta\mathbf{X}_{p,:} \hat{\mathbf{X}}^\top \mathbf{H}_{:,p}^{-1} \right). \quad (6)$$

This reformulation reveals an important property: pruning scores now adaptively update through explicit contributions from previously processed weight columns. Simultaneously, output residual information is implicitly incorporated by fusing the \mathbf{R} updates into Equation 6, allowing \mathbf{R} to evolve across iterations without explicit recalculation.

The next challenge is to compute Equation 6 efficiently. Denoting the three new terms (annotated in red) as $\mathbf{a}_p = \Delta\mathbf{X}_{p,:} \Delta\mathbf{X}_{p,:}^\top$, $\mathbf{b}_p = \Delta\mathbf{X}_{p,:} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} \Delta\mathbf{X}_{p,:}^\top$, and $\mathbf{c}_p = \Delta\mathbf{X}_{p,:} \hat{\mathbf{X}}^\top \mathbf{H}_{:,p}^{-1} / \mathbf{H}_{pp}^{-1}$, efficiently computing Equation 6 is now equivalent of efficiently computing \mathbf{a} , \mathbf{b} , and \mathbf{c} .

First, all elements of \mathbf{a} can be precomputed as $\mathbf{a} = \text{diag}(\Delta\mathbf{X} \Delta\mathbf{X}^\top)$, leveraging the independence across rows of $\Delta\mathbf{X}$. Second, the term \mathbf{b}_p can be reformulated by applying the Cholesky decomposition to \mathbf{H}_{-p}^{-1} , which yields $\mathbf{H}_{-p}^{-1} = \mathbf{L}_{p+1:,p+1:}^\top \mathbf{L}_{p+1:,p+1:}$. Consequently, \mathbf{b}_p transforms into: $\Delta\mathbf{X}_{p,:} \hat{\mathbf{X}}^\top \mathbf{L}_{p+1:,p+1:}^\top \mathbf{L}_{p+1:,p+1:} \hat{\mathbf{X}} \Delta\mathbf{X}_{p,:}^\top$. Exploiting row independence again, the entire \mathbf{b} can be precomputed as $\mathbf{b} = \text{diag}(\mathbf{U} \mathbf{U}^\top)$, where $\mathbf{U}_{p,:} = \Delta\mathbf{X}_{p,:} \hat{\mathbf{X}}^\top \mathbf{L}_{p+1:,p+1:}$.

Finally, by again leveraging the Cholesky decomposition, $\mathbf{H}_{:,p}^{-1} / \mathbf{H}_{pp}^{-1}$ can be easily represented by $\mathbf{L}_{:,p} / \mathbf{L}_{pp}$ (Frantar & Alistarh 2023, Frantar et al. 2022). Thus \mathbf{c} can be precomputed as $\mathbf{c} = \text{diag}(\Delta\mathbf{X} \mathbf{X}^\top \mathbf{L}) \oslash \text{diag}(\mathbf{L})$, where \oslash denotes the element-wise division (Proofs in Appendix).

A further insight is that \mathbf{U} contains components from both the current iteration (the p -th column) and the subsequent one (the $(p+1)$ -th column). Therefore, \mathbf{U} can be re-expressed as: $\mathbf{U} = \left((\Delta\mathbf{X} \mathbf{X}^\top \mathbf{L}) \odot \mathbf{M}^u \right)$, where $\mathbf{M}^u \in \mathcal{R}^{k \times k}$ is a strictly upper-triangular masking matrix with ones above the diagonal.

This reformulation offers a significant computational advantage: the shared intermediate term $\mathbf{Q} = \Delta\mathbf{X} \mathbf{X}^\top \mathbf{L}$ can be reused for computing both \mathbf{b} and \mathbf{c} . Furthermore, \mathbf{U} itself can be reused to precompute the new weight compensation terms in Equation 3, specifically the term $\mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1}$. Denoting this quantity as \mathbf{D} , it can be precomputed as $\mathbf{D} = \mathbf{U} \mathbf{L}^\top$, as demonstrated in Li et al. (2025). The pruning score \mathbf{S} for an arbitrary column p is thus reformed into the following efficient form:

$$\mathbf{S}_{:,p} = \mathbf{W}_{:,p}^2 \left(\frac{1}{\mathbf{H}_{pp}^{-1}} + \mathbf{a}_p - \mathbf{b}_p + 2\mathbf{c}_p \right). \quad (7)$$

Algorithm 1 The **DuoGPT** activation-sparsity-aware pruning for one layer with target of p^v unstructured sparsity. Given lazy batch and mask selection block size B , each consecutive B columns will be p^v sparse.

Input: Dense weight \mathbf{W} , sparse calibration input $\hat{\mathbf{X}}$ with p^x sparsity, Dense input $\tilde{\mathbf{X}}$.
 $\mathbf{H} \leftarrow \hat{\mathbf{X}}\hat{\mathbf{X}}^\top$, $\mathbf{L} = \text{Inverse_Cholesky}(\mathbf{H})$, $\Delta\mathbf{X} \leftarrow \tilde{\mathbf{X}} - \hat{\mathbf{X}}$
 $\mathbf{Q} \leftarrow \Delta\mathbf{X}\hat{\mathbf{X}}^\top\mathbf{L}$, $\mathbf{U} \leftarrow \mathbf{Q} \odot \mathbf{M}_u$
 $\mathbf{a} \leftarrow \text{diag}(\Delta\mathbf{X}\Delta\mathbf{X}^\top)$, $\mathbf{b} \leftarrow \text{diag}(\mathbf{U}\mathbf{U}^\top)$, $\mathbf{c} \leftarrow \text{diag}(\mathbf{Q}) \oslash \text{diag}(\mathbf{L})$, $\mathbf{D} \leftarrow \mathbf{U}\mathbf{L}^\top$
 $\mathbf{P} \leftarrow \mathbf{0}_{n \times k}$, $\mathbf{S} \leftarrow \mathbf{0}_{n \times B}$, $\mathbf{M}^w \leftarrow \mathbf{1}_{n \times B}$, $\mathbf{E} \leftarrow \mathbf{0}_{n \times B}$
for $i = 0, B, 2B, \dots$ **do**
 for $j = i, i + 1, \dots, i + B - 1$ **do**
 if $j \bmod B = 0$ **then**
 $\mathbf{S}_{:,j:j+B} \leftarrow \mathbf{W}_{:,j:j+B}^2 \odot \left(\mathbf{1}^\top \left(\frac{1}{\text{diag}(\mathbf{L})_{j:j+B}^2} + \mathbf{a}_{j:j+B} - \mathbf{b}_{j:j+B} + 2\mathbf{c}_{j:j+B} \right) \right)$
 $\mathbf{M}_{:,j:j+B}^w \leftarrow \text{0-mask } p^v \text{ of weights } w_c \in \mathbf{W}_{:,j:j+B} \text{ with lowest } \mathbf{S}_{:,j:j+B}$
 end if
 $\mathbf{P}_{:,j} \leftarrow \mathbf{W}_{:,j} \odot \mathbf{M}_{:,j}^w$
 $\mathbf{E}_{:,j-i} \leftarrow \frac{(\mathbf{W}_{:,j} - \mathbf{P}_{:,j})}{\mathbf{L}_{jj}}$
 $\mathbf{W}_{:,j:(i+B)} \leftarrow \mathbf{W}_{:,j:(i+B)} - \mathbf{E}_{:,j-i}\mathbf{L}_{j,j:(i+B)}^\top + \mathbf{W}_{:,j}\mathbf{D}_{j,j:(i+B)}$
 end for
 $\mathbf{W}_{:, (i+B):} \leftarrow \mathbf{W}_{:, (i+B):} - \mathbf{E}\mathbf{L}_{i:(i+B), (i+B)}^\top + \mathbf{W}_{:, (i+B):}\mathbf{D}_{i:(i+B), (i+B)}$ // lazy-batch updates
end for
 $\mathbf{W} \leftarrow \mathbf{W} \odot \mathbf{M}^w$

All newly introduced terms are now computable via vectorized operations. Particularly, since both \mathbf{a} and \mathbf{b} are expressed as the diagonal of inner products between matrices, they can be computed using in-place elementwise square operations followed by column-wise summation. As a result, computing the pruning scores for an entire layer reduces to $\mathcal{O}(mk^2)$ complexity—reducing the runtime of naive implementation by a factor of $\min(n, \frac{nk}{m})$.

Moreover, since this efficient reformulation of \mathbf{S} derives directly from the original error metric \mathcal{L} —with the only approximation being the decomposition of \mathbf{R} —the pruning score remains a faithful surrogate to \mathcal{L} , ensuring a close approximation to the exact closed-form solution. The efficient implementation of **DuoGPT** is illustrated in Figure 2(b), with the complete algorithm presented in Algorithm 1. Full derivations are provided in the Appendix.

Compared to the naive implementation with complexity $\mathcal{O}(nmk^2 + nk^3)$, **DuoGPT**’s new complexity is sufficient to make the algorithm practical, even for extremely large models. Additionally, SparseGPT’s iterative blocking and lazy-batch updates are adopted to minimize data movement across iterations, further accelerating runtime. In practice, **DuoGPT** calibrates a LLaMA-3-70B model on a single 80GB A100 GPU in under 140 minutes.

5 Experiments

Experimental Setups. We implement **DuoGPT** using PyTorch (Paszke 2019) and the HuggingFace Transformer library (Wolf et al. 2019) for efficient model and dataset management. All calibration procedures and experiments are performed on 80GB NVIDIA A100 GPUs with offloading (two GPUs are specifically employed for zero-shot evaluations of 70B models). Unless otherwise stated, the calibration dataset consists of 128 2048-token samples, randomly selected from the C4 training dataset (Raffel et al. 2020). To assess model performance, we evaluate the perplexity (PPL) of our **DuoGPT**-pruned models on the WikiText2 dataset (Merity et al. 2016). Furthermore, we complement our evaluation by conducting 0-shot task classifications using the LM Eval Harness (Gao et al. 2021) across widely recognized downstream benchmarks: PIQA (Bisk et al. 2020), HellaSwag (Zellers et al. 2019), WinoGrande (Sakaguchi et al. 2021), ARC-easy, ARC-challenge (Clark et al. 2018), OpenBookQA (OBQA) (Mihaylov et al. 2018), and BoolQ (Clark et al. 2019). More detailed setups can be found in the Appendix.

Unstructured Pruning Baselines Comparisons. We begin by comparing **DuoGPT** with dual-sparse baselines constructed using two widely adopted one-shot unstructured pruning methods: SparseGPT (Frantar & Alistarh 2023) and Wanda (Sun et al. 2023), as shown in Table 1. To ensure a fair comparison, we enable 50% activation sparsity for both **DuoGPT** and SparseGPT during calibration, allowing each method to perform activation-aware compensation. For Wanda, which

Table 1: Comparison against other one-shot unstructured pruning methods in their dual-sparse variants ($\mathbf{W}_{50\%} + \mathbf{X}_{50\%}$) and 2:4 variants ($\mathbf{W}_{2:4} + \mathbf{X}_{\text{dense}}$). **DuoGPT** has 50% dual-sparsity. We also report GPU hours required for calibration. Bold and underlined values indicate the best and second-best results, respectively.

Model	Method	GPU Hours	Wiki2(↓)	PIQA	HellaSwag	ARC-Easy	ARC-Challenge	WinoGrande	Avg(↑)
LLaMA-3-8B	Dense	-	6.14	80.63	79.13	77.53	53.33	72.93	72.71
	SparseGPT	0.21	<u>14.05</u>	<u>72.96</u>	61.97	<u>62.29</u>	37.80	<u>62.51</u>	<u>59.51</u>
	SparseGPT (2:4)	0.18	19.81	70.84	53.39	54.84	32.34	61.56	54.59
	Wanda	0.03	15.98	71.38	57.02	59.39	35.84	61.64	57.05
	Wanda (2:4)	0.03	25.13	67.85	47.97	50.80	29.95	58.96	51.11
	DuoGPT	0.27	13.41	73.72	<u>61.89</u>	63.01	<u>36.77</u>	64.80	60.04
LLaMA-3-70B	Dense	-	2.86	84.60	84.96	86.07	64.25	80.58	80.09
	SparseGPT	1.62	<u>7.54</u>	80.03	74.87	78.49	<u>50.85</u>	<u>74.03</u>	<u>71.65</u>
	SparseGPT (2:4)	1.50	13.32	77.53	63.78	72.64	46.33	72.77	66.61
	Wanda	0.26	8.19	<u>80.03</u>	76.90	77.10	50.09	70.17	70.86
	Wanda (2:4)	0.28	9.32	80.20	73.06	75.42	49.49	71.67	69.97
	DuoGPT	2.28	7.38	80.52	<u>75.94</u>	<u>77.82</u>	52.82	75.69	72.56
LLaMA-2-7B	Dense	-	5.47	79.11	75.99	74.58	46.25	69.06	69.00
	SparseGPT	0.18	<u>8.98</u>	<u>73.88</u>	63.87	<u>62.67</u>	35.75	<u>63.38</u>	<u>60.02</u>
	SparseGPT (2:4)	0.20	12.1	71.65	57.41	59.18	32.17	64.80	57.04
	Wanda	0.02	9.13	72.91	62.33	61.66	35.32	62.59	58.96
	Wanda (2:4)	0.04	12.2	70.95	54.94	57.28	31.40	62.27	55.37
	DuoGPT	0.22	8.58	73.88	<u>63.75</u>	64.18	<u>35.67</u>	65.11	60.52
LLaMA-2-13B	Dense	-	4.88	80.52	79.37	77.53	49.06	72.22	71.74
	SparseGPT	0.32	<u>7.39</u>	<u>76.39</u>	<u>69.21</u>	<u>68.22</u>	<u>41.04</u>	<u>67.64</u>	<u>64.50</u>
	SparseGPT (2:4)	0.33	9.56	74.16	62.95	63.09	37.46	66.22	60.78
	Wanda	0.06	7.41	76.39	69.19	67.51	40.02	65.90	63.80
	Wanda (2:4)	0.07	9.05	75.41	62.69	64.73	37.12	67.09	61.41
	DuoGPT	0.40	7.17	<u>75.63</u>	69.56	69.95	41.13	68.11	64.88
LLaMA-2-70B	Dense	-	3.32	82.75	83.81	81.02	57.34	77.90	76.56
	SparseGPT	1.68	5.09	79.54	77.51	<u>78.87</u>	<u>52.73</u>	75.37	72.80
	SparseGPT (2:4)	1.50	5.94	78.84	74.00	75.34	48.29	75.53	70.40
	Wanda	0.27	<u>5.03</u>	80.52	78.39	77.15	53.16	<u>75.53</u>	<u>72.95</u>
	Wanda (2:4)	0.30	5.47	79.71	76.01	77.36	50.77	75.53	71.88
	DuoGPT	2.30	5.02	<u>80.41</u>	<u>77.64</u>	78.96	52.56	77.51	73.42

lacks weight-adjustment compensation, the calibration data remains in dense format. All methods are evaluated with models pruned to 50% weight sparsity and run with 50% activation sparsity during inference. Additionally, we compare to the semi-unstructured (2:4) variants of these two approaches, which are widely regarded as achieving a balance between speedup and accuracy (Mishra et al. 2021). In the 2:4 setting, activations remain dense during both calibration and evaluation.

Across all LLaMA models, **DuoGPT** consistently achieves the lowest perplexity, validating its capability to enhance pruned model performance under sparse activations. For example, on LLaMA-3-70B, **DuoGPT** reduces perplexity from 7.54 (SparseGPT) to 7.38. In zero-shot evaluations, it surpasses the second-best baseline results on WinoGrande by up to 2.14% and 2.29% for LLaMA-2 and 3, respectively. Overall, **DuoGPT** achieves the highest average accuracy across all evaluated tasks and models. Furthermore, compared to 2:4 weight-only sparse models, **DuoGPT** consistently yields improved performance. Notably, on LLaMA-3-8B, **DuoGPT** enhances the average accuracy from 54.59% (best-performing 2:4 baseline) to 60.04%, marking an improvement of 5.45%. We additionally report GPU hours (single A100 GPU with offloading) required for calibration. Owing to our efficient implementation detailed in Section 4.3, calibration of 70B models can be completed in approximately 2.3 hours.

Table 2: Compare with other structured pruning methods on LLaMA-2-7B. To align the normalized GPU speedup relative to the dense model, we set SliceGPT with 30% and other three baselines with 31.25% weight-only sparsity. **DuoGPT** has 50% dual-sparsity.

Method	Model Size	Speedup	Wiki2(↓)	PIQA	HellaSwag	ARC-Easy	ARC-Challenge	WinoGrande	Avg(↑)
ShortGPT	4.72B/6.74B	1.44×	65.6	63.55	50.92	45.29	<u>33.79</u>	63.22	51.35
2SSP	4.72B/6.74B	1.31×	12.4	<u>71.71</u>	63.11	<u>55.05</u>	31.83	64.17	<u>57.34</u>
BlockPruner	4.72B/6.74B	<u>1.41×</u>	20.5	69.53	54.77	51.89	32.00	60.85	53.81
SliceGPT	5.29B/6.74B	1.26×	23.1	71.22	58.04	56.14	33.11	64.88	56.68
DuoGPT	3.50B/6.74B	1.39×	8.58	73.88	63.75	64.18	35.67	65.11	60.52

Structured Pruning Baselines Comparisons. We compare our dual-sparse **DuoGPT** framework with several popular structured pruning methods: ShortGPT (Men et al. 2024), 2SSP (Sandri et al. 2025), BlockPruner (Zhong et al. 2024), and SliceGPT (Ashkboos et al. 2024a). All methods are evaluated using 256 samples of calibration data (32 for 2SSP, following its original setup). Each model is pruned to approximately 30% structured sparsity, resulting in up to 1.44× end-to-end

speedup over the dense baseline. As shown in Table 2, by leveraging activation sparsity, **DuoGPT** achieves a competitive end-to-end speedup ($1.39\times$) while maintaining significantly better accuracy compared to the structured pruning baselines. In addition, **DuoGPT** yields the smallest model size post-compression, due to its higher 50% unstructured weight sparsity. If structured pruning methods were pushed to a comparable compression ratio of 50%, their performance would degrade sharply.

Table 3: Ablation studies of (a) different dual-sparsity levels and (b) other activation sparse methods.

Method	30%	40%	55%	60%	65%
SparseGPT	5.87	6.52	13.2	29.2	88.2
Wanda	5.85	6.49	16.0	67.1	248
DuoGPT	5.84	6.48	12.3	26.3	77.3

(a) WikiText PPL of pruned LLaMA-2-7B models with varying dual-sparsity.

Method	Model Size	OBQA	Arc-C	BoolQ	Avg
R-Sparse 50%	6.90B/6.74B	30.40	35.49	72.54	46.14
DuoGPT + TEAL 50%	3.50B/6.74B	38.40	36.77	71.77	48.98
DuoGPT 50%	3.50B/6.74B	39.00	35.67	72.48	49.05

(b) Comparison and study of compatibility to other threshold-based activation sparsity work on LLaMA-2-7B models.

Ablation Studies. In Table 3a, we evaluate the performance of our method compared to other dual-sparse baselines across varying levels of dual sparsity. As shown, **DuoGPT** consistently improves the performance of dual-sparse models at all sparsity levels. Notably, the performance gains become more pronounced as the dual sparsity increases.

We further investigate the compatibility of our method with threshold-based activation pruning during evaluation. As shown in Table 3b, applying the activation thresholding technique from TEAL (Liu et al. 2025) to a 50% dual-sparse **DuoGPT** model yields comparable performance to on-the-fly magnitude-based pruning. We also compare against R-Sparse (Zhang et al. 2025), a prior approach that combines activation sparsity with SVD-based low-rank weight compression. As Table 3b shows, **DuoGPT** achieves higher average accuracy than R-Sparse under 50% sparsity. Moreover, while our method reduces the model size to 50% of the dense baseline, R-Sparse incurs a 1% increase in memory usage due to the additional SVD branch.




Figure 3: Number of weights of LLaMA-2-7B loaded to SRAM. SWS=structured weight sparsity, AS=activation sparsity, and DS=dual sparsity.

Next, we analyze the number of weights loaded into SRAM across different sparsity paradigms, as shown in Figure 3. Consistent with our discussion in Section 4.1, we observe that the dual-sparse model loads approximately $(1-p^w) \times (1-p^x)$ fraction of the weights into SRAM. This result shows potential for further improvement of the efficiency of dual-sparse LLMs with custom GPU spMSPV kernels or ASIC dual-sparse accelerators.

Table 4: Comparison with other baselines on 5-shot GSM8K with LLaMA-3-8B.

Method (X, W) _{sparsity}	Model Size	SRAM Loads	GSM8K
TEAL (50%, dense)	8.03B/8.03B	50.0%	36.47
SparseGPT (30%, 30%)	5.94B/8.03B	49.7%	36.39
DuoGPT (30%, 30%)	5.94B/8.03B	49.6%	36.77

Finally, we evaluate the performance of **DuoGPT** on a reasoning benchmark. Table 4 compares the accuracy of 50% activation-sparse TEAL, 30% dual-sparse SparseGPT, and **DuoGPT** on the 5-shot GSM8K (Cobbe et al. 2021) using LLaMA-3-8B, under the same number of SRAM weight fetches. **DuoGPT** achieves the highest accuracy and the smallest model size under iso-SRAM weight load.

Limitations and Future Work. This work does not yet include a full GPU kernel implementation for accelerating the spMSPV operation. Unlike spMSPM, spMSPV avoids costly index matching, making it a strong candidate for efficient GPU acceleration—an avenue we leave for future exploration. While **DuoGPT** offers a promising approach to calibrate dual-sparse LLMs, this remains an open problem with substantial potential for further research.

6 Conclusion

In this paper, we present **DuoGPT**, an efficient one-shot dual-sparse pruning framework that unifies unstructured weight pruning with runtime activation sparsity. By extending the OBC framework with activation-aware calibration and residual correction, **DuoGPT** achieves state-of-the-art accuracy while significantly improving both memory and computational efficiency. Our method scales to large LLMs and complements existing activation pruning techniques, offering a practical path toward efficient and scalable LLM deployment.

References

- Ashkboos, S., Croci, M. L., Nascimento, M. G. d., Hoefler, T., and Hensman, J. Slicept: Compress large language models by deleting rows and columns. *arXiv preprint arXiv:2401.15024*, 2024a.
- Ashkboos, S., Mohtashami, A., Croci, M., Li, B., Cameron, P., Jaggi, M., Alistarh, D., Hoefler, T., and Hensman, J. Quarot: Outlier-free 4-bit inference in rotated llms. *Advances in Neural Information Processing Systems*, 37:100213–100240, 2024b.
- Bisk, Y., Zellers, R., Gao, J., Choi, Y., et al. Piqa: Reasoning about physical commonsense in natural language. In *Proceedings of the AAAI conference on artificial intelligence*, volume 34, pp. 7432–7439, 2020.
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- Clark, C., Lee, K., Chang, M.-W., Kwiatkowski, T., Collins, M., and Toutanova, K. Boolq: Exploring the surprising difficulty of natural yes/no questions. *arXiv preprint arXiv:1905.10044*, 2019.
- Clark, P., Cowhey, I., Etzioni, O., Khot, T., Sabharwal, A., Schoenick, C., and Tafjord, O. Think you have solved question answering? try arc, the ai2 reasoning challenge. *arXiv preprint arXiv:1803.05457*, 2018.
- Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Frantar, E. and Alistarh, D. Optimal brain compression: A framework for accurate post-training quantization and pruning. *Advances in Neural Information Processing Systems*, 35:4475–4488, 2022.
- Frantar, E. and Alistarh, D. Sparsegpt: Massive language models can be accurately pruned in one-shot. In *International Conference on Machine Learning*, pp. 10323–10337. PMLR, 2023.
- Frantar, E., Ashkboos, S., Hoefler, T., and Alistarh, D. Gptq: Accurate post-training quantization for generative pre-trained transformers. *arXiv preprint arXiv:2210.17323*, 2022.
- Gao, L., Tow, J., Biderman, S., Black, S., DiPofi, A., Foster, C., Golding, L., Hsu, J., McDonell, K., Muennighoff, N., et al. A framework for few-shot language model evaluation. *Version v0. 0.1. Sept*, 10:8–9, 2021.
- Han, S., Pool, J., Tran, J., and Dally, W. Learning both weights and connections for efficient neural network. *Advances in neural information processing systems*, 28, 2015.
- Hassibi, B. and Stork, D. Second order derivatives for network pruning: Optimal brain surgeon. *Advances in neural information processing systems*, 5, 1992.
- Haziza, D., Chou, T., Choudhary, D., Wehrstedt, L., Massa, F., Yu, J., Jeong, G., Rao, S., Labatut, P., and Cai, J. Accelerating transformer inference and training with 2: 4 activation sparsity. *arXiv preprint arXiv:2503.16672*, 2025.
- LeCun, Y., Denker, J., and Solla, S. Optimal brain damage. *Advances in neural information processing systems*, 2, 1989.
- Lee, D., Lee, J., Zhang, G., Tiwari, M., and Mirhoseini, A. Cats: Context-aware thresholding for sparsity in large language models. In *First Conference on Language Modeling*, 2024.
- Li, Y., Yin, R., Lee, D., Xiao, S., and Panda, P. Gptqv2: Efficient finetuning-free quantization for asymmetric calibration. *arXiv preprint arXiv:2504.02692*, 2025.
- Li, Z., You, C., Bhojanapalli, S., Li, D., Rawat, A. S., Reddi, S. J., Ye, K., Chern, F., Yu, F., Guo, R., et al. The lazy neuron phenomenon: On emergence of activation sparsity in transformers. *arXiv preprint arXiv:2210.06313*, 2022.

- Lin, C.-H., Gao, S., Smith, J. S., Patel, A., Tuli, S., Shen, Y., Jin, H., and Hsu, Y.-C. Modegpt: Modular decomposition for large language model compression. *arXiv preprint arXiv:2408.09632*, 2024.
- Ling, G., Wang, Z., and Liu, Q. Slimgpt: Layer-wise structured pruning for large language models. *Advances in Neural Information Processing Systems*, 37:107112–107137, 2024.
- Liu, J., Ponnusamy, P., Cai, T., Guo, H., Kim, Y., and Athiwaratkun, B. Training-free activation sparsity in large language models. *The Thirteenth International Conference on Learning Representations*, 2025.
- Liu, Z., Wang, J., Dao, T., Zhou, T., Yuan, B., Song, Z., Shrivastava, A., Zhang, C., Tian, Y., Re, C., et al. Deja vu: Contextual sparsity for efficient llms at inference time. In *International Conference on Machine Learning*, pp. 22137–22176. PMLR, 2023.
- Men, X., Xu, M., Zhang, Q., Wang, B., Lin, H., Lu, Y., Han, X., and Chen, W. Shortgpt: Layers in large language models are more redundant than you expect. *arXiv preprint arXiv:2403.03853*, 2024.
- Merity, S., Xiong, C., Bradbury, J., and Socher, R. Pointer sentinel mixture models. *arXiv preprint arXiv:1609.07843*, 2016.
- Mihaylov, T., Clark, P., Khot, T., and Sabharwal, A. Can a suit of armor conduct electricity? a new dataset for open book question answering. *arXiv preprint arXiv:1809.02789*, 2018.
- Mirzadeh, I., Alizadeh, K., Mehta, S., Del Mundo, C. C., Tuzel, O., Samei, G., Rastegari, M., and Farajtabar, M. Relu strikes back: Exploiting activation sparsity in large language models. *arXiv preprint arXiv:2310.04564*, 2023.
- Mishra, A., Latorre, J. A., Pool, J., Stosic, D., Stosic, D., Venkatesh, G., Yu, C., and Micikevicius, P. Accelerating sparse deep neural networks. *arXiv preprint arXiv:2104.08378*, 2021.
- Paszke, A. Pytorch: An imperative style, high-performance deep learning library. *arXiv preprint arXiv:1912.01703*, 2019.
- Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., and Liu, P. J. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of machine learning research*, 21(140):1–67, 2020.
- Sakaguchi, K., Bras, R. L., Bhagavatula, C., and Choi, Y. Winogrande: An adversarial winograd schema challenge at scale. *Communications of the ACM*, 64(9):99–106, 2021.
- Sandri, F., Cunegatti, E., and Iacca, G. 2ssp: A two-stage framework for structured pruning of llms. *arXiv preprint arXiv:2501.17771*, 2025.
- Sun, M., Liu, Z., Bair, A., and Kolter, J. Z. A simple and effective pruning approach for large language models. *arXiv preprint arXiv:2306.11695*, 2023.
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Wang, X., Zheng, Y., Wan, Z., and Zhang, M. Svd-llm: Truncation-aware singular value decomposition for large language model compression. *arXiv preprint arXiv:2403.07378*, 2024.
- Wolf, T., Debut, L., Sanh, V., Chaumond, J., Delangue, C., Moi, A., Cistac, P., Rault, T., Louf, R., Funtowicz, M., et al. Huggingface’s transformers: State-of-the-art natural language processing. *arXiv preprint arXiv:1910.03771*, 2019.
- Zellers, R., Holtzman, A., Bisk, Y., Farhadi, A., and Choi, Y. Hellaswag: Can a machine really finish your sentence? *arXiv preprint arXiv:1905.07830*, 2019.
- Zhang, Z., Lin, Y., Liu, Z., Li, P., Sun, M., and Zhou, J. Moefication: Transformer feed-forward layers are mixtures of experts. *arXiv preprint arXiv:2110.01786*, 2021.

- Zhang, Z., Song, Y., Yu, G., Han, X., Lin, Y., Xiao, C., Song, C., Liu, Z., Mi, Z., and Sun, M. Relu² wins: Discovering efficient activation functions for sparse llms. *arXiv preprint arXiv:2402.03804*, 2024a.
- Zhang, Z., Xiao, C., Qin, Q., Lin, Y., Zeng, Z., Han, X., Liu, Z., Xie, R., Sun, M., and Zhou, J. Exploring the benefit of activation sparsity in pre-training. *arXiv preprint arXiv:2410.03440*, 2024b.
- Zhang, Z., Liu, Z., Tian, Y., Khaitan, H., Wang, Z., and Li, S. R-sparse: Rank-aware activation sparsity for efficient llm inference. In *The Thirteenth International Conference on Learning Representations*, 2025.
- Zhong, L., Wan, F., Chen, R., Quan, X., and Li, L. Blockpruner: Fine-grained pruning for large language models. *arXiv preprint arXiv:2406.10594*, 2024.

A Theoretical Derivation

A.1 Proof of Equation 3 (Optimal Framework for Activation Sparsity aware Pruning)

We provide the detailed derivation of **DuoGPT**'s optimal calibration framework (Equation 3) proposed in Section 4.2.

Proof. The Lagrangian in Equation 2 is first expanded:

$$\begin{aligned}\mathcal{L} &= \text{Tr}\left((\Delta\mathbf{w}\hat{\mathbf{X}} - \mathbf{r})^\top (\Delta\mathbf{w}\hat{\mathbf{X}} - \mathbf{r})\right) + \lambda\Delta\mathbf{w}\mathbf{e}_p^\top + \lambda\mathbf{w}_p \\ &= \text{Tr}(\hat{\mathbf{X}}^\top \Delta\mathbf{w}^\top \Delta\mathbf{w}\hat{\mathbf{X}} - \hat{\mathbf{X}}^\top \Delta\mathbf{w}^\top \mathbf{r} - \mathbf{r}^\top \Delta\mathbf{w}\hat{\mathbf{X}} + \mathbf{r}^\top \mathbf{r}) + \lambda\Delta\mathbf{w}\mathbf{e}_p^\top + \lambda\mathbf{w}_p\end{aligned}\quad (8)$$

To find the local minima of the Lagrangian in Equation 2, we set the partial derivatives with respect to $\Delta\mathbf{w}$ and λ to zero. We first compute the partial derivative with respect to $\Delta\mathbf{w}$:

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial \Delta\mathbf{w}} &= 2\Delta\mathbf{w}\hat{\mathbf{X}}\hat{\mathbf{X}}^\top - 2\mathbf{r}\hat{\mathbf{X}}^\top + \lambda\mathbf{e}_p \\ &= 2\Delta\mathbf{w}\mathbf{H} - 2\mathbf{r}\hat{\mathbf{X}}^\top + \lambda\mathbf{e}_p,\end{aligned}\quad (9)$$

and the partial derivative with respect to λ :

$$\frac{\partial \mathcal{L}}{\partial \lambda} = \Delta\mathbf{w}\mathbf{e}_p^\top + \mathbf{w}_p. \quad (10)$$

Setting the first equation to zero yields:

$$\begin{aligned}\Delta\mathbf{w} &= \mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}^{-1} - \frac{\lambda}{2}\mathbf{e}_p \mathbf{H}^{-1} \\ &= \mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}^{-1} - \frac{\lambda}{2}\mathbf{H}_{p,:}^{-1}.\end{aligned}\quad (11)$$

Setting Equation 10 to zero and substituting $\Delta\mathbf{w}$ from Equation 11, we obtain:

$$\mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}^{-1} \mathbf{e}_p^\top - \frac{\lambda}{2}\mathbf{H}_{p,:}^{-1} \mathbf{e}_p^\top + \mathbf{w}_p = 0. \quad (12)$$

By simplifying with $\mathbf{H}_{p,:}^{-1} \mathbf{e}_p^\top = \mathbf{H}_{pp}^{-1}$, we solve for λ as:

$$\begin{aligned}\lambda &= \frac{2}{\mathbf{H}_{pp}^{-1}} (\mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}^{-1} \mathbf{e}_p^\top + \mathbf{w}_p) \\ &= \frac{2}{\mathbf{H}_{pp}^{-1}} (\mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}_{:,p}^{-1} + \mathbf{w}_p)\end{aligned}\quad (13)$$

Substituting Equation 13 back into Equation 11, we solve for $\Delta\mathbf{w}$ as:

$$\begin{aligned}\Delta\mathbf{w} &= \mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}^{-1} - \frac{1}{\mathbf{H}_{pp}^{-1}} (\mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}_{:,p}^{-1} \mathbf{H}_{p,:}^{-1} + \mathbf{w}_p \mathbf{H}_{p,:}^{-1}) \\ &= -\frac{\mathbf{w}_p}{\mathbf{H}_{pp}^{-1}} \cdot \mathbf{H}_{p,:}^{-1} + \mathbf{r}\hat{\mathbf{X}}^\top (\mathbf{H}^{-1} - \frac{\mathbf{H}_{:,p}^{-1} \mathbf{H}_{p,:}^{-1}}{\mathbf{H}_{pp}^{-1}}).\end{aligned}\quad (14)$$

By recognizing that the expression enclosed in parentheses in the second term corresponds to the Gaussian elimination operation introduced in Section 3, we obtain the final expression for $\Delta\mathbf{w}$:

$$\Delta\mathbf{w} = -\frac{\mathbf{w}_p}{\mathbf{H}_{pp}^{-1}} \cdot \mathbf{H}_{p,:}^{-1} + \mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1}, \quad (15)$$

which is identical to the weight update term in Equation 3.

The next step is to derive the final expression for \mathcal{L} . Substituting Equation 15 into the loss function $\mathcal{L} = \|\Delta\mathbf{w}\hat{\mathbf{X}} - \mathbf{r}\|_F^2$, the loss becomes:

$$\mathcal{L} = \left\| -\frac{\mathbf{w}_p}{\mathbf{H}_{pp}^{-1}} \mathbf{H}_{p,:}^{-1} \hat{\mathbf{X}} + \mathbf{r}\hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} - \mathbf{r} \right\|_F^2. \quad (16)$$

This expression can be expanded as:

$$\begin{aligned}
\mathcal{L} &= \frac{\mathbf{w}_p^2}{(\mathbf{H}_{pp}^{-1})^2} \mathbf{H}_{p,:}^{-1} \mathbf{H} \mathbf{H}_{:,p}^{-1} - 2(\mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} - \mathbf{r}) \left(\frac{\mathbf{w}_p}{\mathbf{H}_{pp}^{-1}} \hat{\mathbf{X}}^\top \mathbf{H}_{:,p}^{-1} \right) + \mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \mathbf{H} \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} \mathbf{r}^\top \\
&\quad - 2\mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} \mathbf{r}^\top + \mathbf{r} \mathbf{r}^\top \\
&= \frac{\mathbf{w}_p^2}{(\mathbf{H}_{pp}^{-1})^2} \mathbf{H}_{p,:}^{-1} \mathbf{H} \mathbf{H}_{:,p}^{-1} - 2 \frac{\mathbf{w}_p}{\mathbf{H}_{pp}^{-1}} \mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \mathbf{H} \mathbf{H}_{:,p}^{-1} + 2 \frac{\mathbf{w}_p}{\mathbf{H}_{pp}^{-1}} \mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{:,p}^{-1} \\
&\quad + \mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \mathbf{H} \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} \mathbf{r}^\top - 2\mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} \mathbf{r}^\top + \mathbf{r} \mathbf{r}^\top.
\end{aligned} \tag{17}$$

To simplify Equation 17, we observe that the second and third terms are nearly identical, differing only by the presence of the term $\mathbf{H}_{-p}^{-1} \mathbf{H}$. This same pattern appears between the fourth and fifth terms. We begin by simplifying $\mathbf{H}_{-p}^{-1} \mathbf{H}$:

$$\begin{aligned}
\mathbf{H}_{-p}^{-1} \mathbf{H} &= \left(\mathbf{H}^{-1} - \frac{\mathbf{H}_{:,p}^{-1} \mathbf{H}_{p,:}^{-1}}{\mathbf{H}_{pp}^{-1}} \right) \mathbf{H} \\
&= \mathbf{I} - \frac{\mathbf{H}^{-1} \mathbf{e}_p^\top \mathbf{e}_p \mathbf{H}^{-1}}{\mathbf{H}_{pp}^{-1}} \mathbf{H} \\
&= \mathbf{I} - \frac{1}{\mathbf{H}_{pp}^{-1}} \mathbf{H}^{-1} \mathbf{e}_p^\top \mathbf{e}_p.
\end{aligned} \tag{18}$$

With the help of the identity matrix, we can more clearly see how the aforementioned nearly identical terms simplify. Substituting Equation 18 back into Equation 17:

$$\begin{aligned}
\mathcal{L} &= \frac{\mathbf{w}_p^2}{(\mathbf{H}_{pp}^{-1})^2} \mathbf{H}_{p,:}^{-1} \mathbf{H} \mathbf{H}_{:,p}^{-1} + 2 \frac{\mathbf{w}_p}{(\mathbf{H}_{pp}^{-1})^2} \mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}^{-1} \mathbf{e}_p^\top \mathbf{e}_p \mathbf{H}_{:,p}^{-1} \\
&\quad - \frac{1}{\mathbf{H}_{pp}^{-1}} \mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}^{-1} \mathbf{e}_p^\top \mathbf{e}_p \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} \mathbf{r}^\top - \mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} \mathbf{r}^\top + \mathbf{r} \mathbf{r}^\top.
\end{aligned} \tag{19}$$

An important observation is that $\mathbf{e}_p \mathbf{H}_{-p}^{-1}$ is in fact an all-zero vector since the p -th row of \mathbf{H}_{-p}^{-1} is eliminated. Consequently, the third term in Equation 19 equals zero. Part of the second term can also be further simplified: $\mathbf{H}^{-1} \mathbf{e}_p^\top \mathbf{e}_p \mathbf{H}_{:,p}^{-1} = \mathbf{H}_{:,p}^{-1} \mathbf{H}_{pp}^{-1}$. The expression $\mathbf{H}_{p,:}^{-1} \mathbf{H} \mathbf{H}_{:,p}^{-1}$ in the first term can also be simplified to \mathbf{H}_{pp}^{-1} . Combining all these simplifications, Equation 19 reduces to:

$$\mathcal{L} = \frac{\mathbf{w}_p^2}{\mathbf{H}_{pp}^{-1}} + 2 \frac{\mathbf{w}_p}{\mathbf{H}_{pp}^{-1}} \mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{:,p}^{-1} - \mathbf{r} \hat{\mathbf{X}}^\top \mathbf{H}_{-p}^{-1} \hat{\mathbf{X}} \mathbf{r}^\top + \mathbf{r} \mathbf{r}^\top. \tag{20}$$

By reordering the second and fourth terms, we obtain the identical format of \mathcal{L} as presented in Equation 3.

A.2 Proof of Cholesky Decomposition of \mathbf{H}_{-p}^{-1} (Efficient calculation of \mathbf{b} for pruning score \mathbf{S})

Recall that in Section 4.3, we leverage the equality of $\mathbf{H}_{-p}^{-1} = \mathbf{L}_{p+1:,p+1:} \mathbf{L}_{p+1:,p+1:}^\top$ to simplify the computation of term \mathbf{b} for pruning score \mathbf{S} . We provide the proof for this equality using mathematical induction below.

Proof. We begin the mathematical induction proof by establishing the base case: $p = 1$.

Base Case: First, we rewrite the lower-triangular Cholesky factor \mathbf{L} as:

$$\mathbf{L} = \begin{pmatrix} \mathbf{L}_{11} & \mathbf{0} \\ \mathbf{L}_{2:,1} & \mathbf{L}_{2:,2:} \end{pmatrix}. \tag{21}$$

The Cholesky decomposition of \mathbf{H} can be written as:

$$\mathbf{H}^{-1} = \begin{pmatrix} \mathbf{H}_{11}^{-1} & \mathbf{H}_{1,2:}^{-1} \\ \mathbf{H}_{2:,1}^{-1} & \mathbf{H}_{2:,2:}^{-1} \end{pmatrix} = \begin{pmatrix} \mathbf{L}_{11} & \mathbf{0} \\ \mathbf{L}_{2:,1} & \mathbf{L}_{2:,2:} \end{pmatrix} \begin{pmatrix} \mathbf{L}_{11} & \mathbf{L}_{2:,1}^\top \\ \mathbf{0} & \mathbf{L}_{2:,2:}^\top \end{pmatrix}. \tag{22}$$

Based on this equation, we can construct the following linear system:

$$\begin{cases} \mathbf{H}_{11}^{-1} = \mathbf{L}_{11}^2 \\ \mathbf{H}_{2:,1}^{-1} = \mathbf{L}_{11}\mathbf{L}_{2:,1} \\ \mathbf{H}_{2:,2}^{-1} = \mathbf{L}_{2:,1}\mathbf{L}_{2:,1}^\top + \mathbf{L}_{2:,2}\mathbf{L}_{2:,2}^\top \end{cases}. \quad (23)$$

Solving the above equations, we obtain the expression for $\mathbf{L}_{2:,2}\mathbf{L}_{2:,2}^\top$ as:

$$\begin{aligned} \mathbf{L}_{2:,2}\mathbf{L}_{2:,2}^\top &= \mathbf{H}_{2:,2}^{-1} - \frac{1}{\sqrt{\mathbf{H}_{11}^{-1}}}\mathbf{H}_{2:,1}^{-1}\frac{1}{\sqrt{\mathbf{H}_{11}^{-1}}}(\mathbf{H}_{2:,1}^{-1})^\top \\ &= \mathbf{H}_{2:,2}^{-1} - \frac{1}{\mathbf{H}_{11}^{-1}}\mathbf{H}_{2:,1}^{-1}\mathbf{H}_{1,2}^{-1}. \end{aligned} \quad (24)$$

Recalling the expression for Gaussian elimination introduced in Section 3, Equation 24 essentially represents the removal of the first row and column from \mathbf{H}^{-1} after applying Gaussian elimination to zero out its first row and column. Therefore, we have proven the base case that $\mathbf{H}_{-1}^{-1} = \mathbf{L}_{2:,2}\mathbf{L}_{2:,2}^\top$.

Inductive Step: Assume the statement holds for $p = k - 1$. We show it also holds for $p = k$. Since we know $\mathbf{H}_{-(k-1)}^{-1} = \mathbf{L}_{k:,k}\mathbf{L}_{k:,k}^\top$, the lower-triangular Cholesky factor \mathbf{L} for $\mathbf{H}_{-(k-1)}^{-1}$ can be formulated as:

$$\mathbf{L} = \begin{pmatrix} \mathbf{L}_{kk} & \mathbf{0} \\ \mathbf{L}_{(k+1):,k} & \mathbf{L}_{(k+1):,(k+1):} \end{pmatrix}. \quad (25)$$

This formulation is valid because, under our inductive hypothesis for $p = k - 1$, all rows and columns before k have already been removed. We can construct similar linear systems as in the base case:

$$\begin{cases} \mathbf{H}_{kk}^{-1} = \mathbf{L}_{kk}^2 \\ \mathbf{H}_{(k+1):,k}^{-1} = \mathbf{L}_{kk}\mathbf{L}_{(k+1):,k} \\ \mathbf{H}_{(k+1):,(k+1):}^{-1} = \mathbf{L}_{(k+1):,k}\mathbf{L}_{(k+1):,k}^\top + \mathbf{L}_{(k+1):,(k+1):}\mathbf{L}_{(k+1):,(k+1):}^\top \end{cases}. \quad (26)$$

Solving the above equations, we again obtain:

$$\begin{aligned} \mathbf{L}_{(k+1):,(k+1):}\mathbf{L}_{(k+1):,(k+1):}^\top &= \mathbf{H}_{(k+1):,(k+1):}^{-1} - \frac{1}{\sqrt{\mathbf{H}_{kk}^{-1}}}\mathbf{H}_{(k+1):,k}^{-1}\frac{1}{\sqrt{\mathbf{H}_{kk}^{-1}}}(\mathbf{H}_{(k+1):,k}^{-1})^\top \\ &= \mathbf{H}_{(k+1):,(k+1):}^{-1} - \frac{1}{\mathbf{H}_{kk}^{-1}}\mathbf{H}_{(k+1):,k}^{-1}\mathbf{H}_{k,(k+1):}^{-1}. \end{aligned} \quad (27)$$

Again, Equation 27 is equivalent to removing the first row and column of $\mathbf{H}_{-(k-1)}^{-1}$ after applying the Gaussian Elimination to zero out its first row and column. This operation is equivalent to computing \mathbf{H}_{-k}^{-1} . Therefore, we have proven that $\mathbf{H}_{-k}^{-1} = \mathbf{L}_{(k+1):,(k+1):}\mathbf{L}_{(k+1):,(k+1):}^\top$. By the principle of mathematical induction, the statement holds for all p with $k \geq 1$.

A.3 Proof of $\mathbf{H}_{:,p}^{-1}/\mathbf{H}_{pp}^{-1} = \mathbf{L}_{:,p}/\mathbf{L}_{pp}$ (Efficient calculation of \mathbf{c} for pruning score \mathbf{S})

This equality has previously been leveraged to simplify the implementation of SparseGPT (Frantar & Alistarh 2023) and GPTQ (Frantar et al. 2022). We also leverage this relationship to simplify the calculation of term \mathbf{c} in our pruning score \mathbf{S} .

Proof. We assume that at the p -th iteration of the calibration, all columns and rows before p have been removed through Gaussian elimination. Based on this assumption, the Cholesky decomposition of the Hessian inverse at the p -th iteration can be formulated as:

$$\mathbf{H}_{-(p-1)}^{-1} = \begin{pmatrix} \mathbf{H}_{pp}^{-1} & \mathbf{H}_{p,(p+1):}^{-1} \\ \mathbf{H}_{(p+1):,p}^{-1} & \mathbf{H}_{(p+1):,(p+1):}^{-1} \end{pmatrix} = \begin{pmatrix} \mathbf{L}_{pp} & \mathbf{0} \\ \mathbf{L}_{(p+1):,p} & \mathbf{L}_{(p+1):,(p+1):} \end{pmatrix} \begin{pmatrix} \mathbf{L}_{pp} & \mathbf{L}_{(p+1):,p}^\top \\ \mathbf{0} & \mathbf{L}_{(p+1):,(p+1):}^\top \end{pmatrix}. \quad (28)$$

This equation also leverages the result proven in Section A.2. Next, we divide the expression $\mathbf{H}_{:,p}^{-1}$ into three parts: $\mathbf{H}_{1:(p-1),p}^{-1}$, \mathbf{H}_{pp}^{-1} , and $\mathbf{H}_{(p+1):,p}^{-1}$. For the first two parts, the equality is straightforward

to prove. Since $\mathbf{H}_{1:(p-1),p}^{-1} = \mathbf{0}^\top = \mathbf{L}_{1:(p-1),p}$, so $\mathbf{H}_{1:(p-1),p}^{-1}/\mathbf{H}_{pp}^{-1} = \mathbf{0}^\top = \mathbf{L}_{1:(p-1),p}/\mathbf{L}_{pp}$. And $\mathbf{H}_{pp}^{-1}/\mathbf{H}_{pp}^{-1} = 1 = \mathbf{L}_{pp}/\mathbf{L}_{pp}$. The only equality that requires proof is: $\mathbf{H}_{(p+1):,p}^{-1}/\mathbf{H}_{pp}^{-1} = \mathbf{L}_{(p+1):,p}/\mathbf{L}_{pp}$. From Equation 28, we can construct the following linear relations:

$$\begin{cases} \mathbf{H}_{pp}^{-1} = \mathbf{L}_{pp}^2 \\ \mathbf{H}_{(p+1):,p}^{-1} = \mathbf{L}_{pp}\mathbf{L}_{(p+1):,p} \end{cases}. \quad (29)$$

Substituting these relations back into the expression, it is straightforward to verify that $\mathbf{H}_{(p+1):,p}^{-1}/\mathbf{H}_{pp}^{-1} = \mathbf{L}_{pp}\mathbf{L}_{(p+1):,p}/\mathbf{L}_{pp}^2 = \mathbf{L}_{(p+1):,p}/\mathbf{L}_{pp}$.

A.4 Proof of $\mathbf{U} = ((\Delta\mathbf{X}\mathbf{X}^\top\mathbf{L}) \odot \mathbf{M}^u)$

Recall that in Section 4.2, \mathbf{b} is first proposed to be precomputed as $\mathbf{b} = \text{diag}(\mathbf{U}\mathbf{U}^\top)$, where $\mathbf{U}_{p,:} = \Delta\mathbf{X}_{p,:}\hat{\mathbf{X}}^\top\mathbf{L}_{p+1:,p+1:}$. To further improve the algorithm efficiency, we propose to re-express \mathbf{U} in the form $((\Delta\mathbf{X}\mathbf{X}^\top\mathbf{L}) \odot \mathbf{M}^u)$, so that the term $\Delta\mathbf{X}\mathbf{X}^\top\mathbf{L}$ can be reused between \mathbf{b} and \mathbf{c} . The proof is straightforward and provided below.

Proof. The p -th row of \mathbf{U} equals $\Delta\mathbf{X}_{p,:}\hat{\mathbf{X}}^\top\mathbf{L}_{p+1:,p+1:}$, where $\mathbf{L} \in \mathcal{R}^{k \times k}$ is the lower-triangle Cholesky factor. For any given vector $\mathbf{z} \in \mathcal{R}^{1 \times k}$ and the Cholesky lower-triangle factor matrix $\mathbf{L}_{p+1:,p+1:}$, we have:

$$(\mathbf{z}\mathbf{L}_{p+1:,p+1:})_i = \begin{cases} 0 & \text{if } i < p+1 \\ \mathbf{z}\mathbf{L}_{:,i} & \text{else } i \geq p+1 \end{cases}. \quad (30)$$

Therefore, substituting $\mathbf{z} = \Delta\mathbf{X}_{p,:}\hat{\mathbf{X}}^\top \in \mathcal{R}^{1 \times k}$ into the above equation, we can reformulate the p -th row of \mathbf{U} as:

$$\mathbf{U}_{p,i} = \begin{cases} 0 & \text{if } i < p+1 \\ \Delta\mathbf{X}_{p,:}\hat{\mathbf{X}}^\top\mathbf{L}_{:,i} & \text{else } i \geq p+1 \end{cases}. \quad (31)$$

Hence, $\mathbf{U}_{p,:}$ equals $(\Delta\mathbf{X}_{p,:}\hat{\mathbf{X}}^\top\mathbf{L}) \odot \mathbf{M}_{p,:}^u$, where \mathbf{M}^u is a strictly upper-triangular masking matrix with ones above the diagonal. The full computation of \mathbf{U} can thus be re-expressed as $\mathbf{U} = (\Delta\mathbf{X}\hat{\mathbf{X}}^\top\mathbf{L}) \odot \mathbf{M}^u$. By leveraging the independence across rows of $\Delta\mathbf{X}$, $\mathbf{b} = \text{diag}(\mathbf{U}\mathbf{U}^\top)$, where $\mathbf{U} = (\Delta\mathbf{X}\hat{\mathbf{X}}^\top\mathbf{L}) \odot \mathbf{M}^u$.

B Detailed Experimental Setups

This section provides comprehensive implementation details and experimental configurations to ensure reproducibility of our results.

B.1 Calibration Settings for DuoGPT, and Unstructured Baselines

We implement all unstructured baselines reported in Table 1 using uniform evaluation settings to ensure fair comparison. Our implementation builds upon the calibration framework from SparseGPT (Frantar & Alistarh 2023), using PyTorch (Paszke 2019).

For both SparseGPT and **DuoGPT**, we enable *act_order*, an option in SparseGPT that sorts weight columns based on Hessian diagonal magnitude to improve pruning performance. The dampening ratio for Hessian is set to 0.1 for numerical stability. We apply the lazy batch updates and iterative mask blocking techniques from SparseGPT to **DuoGPT** for reducing IO overhead. The block size and lazy batch size are both set to $B = 128$. For both methods, we prune the input calibration data to each layer based on magnitude before calculating each layer’s Hessians.

We integrate the original pruning and mask selection procedures from Wanda (Sun et al. 2023) into our evaluation framework. Since Wanda lacks a compensation mechanism post-pruning, we maintain dense activation throughout the calibration process (activation sparsity aware pruning disabled).

For 2:4 structured versions of both SparseGPT and Wanda, we apply mask selection and pruning over every 4-column group with 50% weight sparsity. All other configurations remain identical to their unstructured counterparts. Calibration data remains dense throughout the pruning process for all 2:4 baselines.

B.2 Settings for Speedup Results

We report end-to-end GPU speedup results on an NVIDIA A100 80GB GPU in Table 2 and 5. All speedups are normalized with respect to the dense model baseline rather than reporting absolute running times to ensure fair comparison across different frameworks.

For **DuoGPT**, end-to-end GPU speedup is measured by inheriting TEAL’s custom Triton kernel (Liu et al. 2025) and integrating it into the fast-gpt framework. The corresponding dense model baseline is also evaluated using the fast-gpt framework to obtain normalized speedup ratios. All speedups for structured pruned models are measured using 2SSP’s framework (Sandri et al. 2025), with dense model baselines similarly evaluated within the same framework to compute normalized speedup.

B.3 Settings to Measure the Number of Weights Loaded to SRAM and Model Size

In Figure 3, we report the number of weights loaded into SRAM during inference. For dense models, structured-sparse models, and activation-sparse models, the number of weights loaded into SRAM is fixed. However, for **DuoGPT**, the dual-sparsity approach combined with unstructured weight distribution across rows makes the number of loaded weights a dynamic quantity that depends on which activation patterns are selected. To ensure conservative and fair comparison, we report worst-case SRAM loading numbers for **DuoGPT**. Specifically, we sort weight rows by their weight sparsity and assume that dynamic activation sparsity always selects the rows with the lowest weight sparsity (i.e., the densest rows requiring maximum SRAM access). For example, with 40% activation sparsity and 40% unstructured weight sparsity, we first sort all weight rows by their sparsity levels, then select the 40% of rows with the lowest weight sparsity—those requiring the largest number of weights to be loaded into SRAM. This conservative methodology ensures our reported SRAM access numbers represent an upper bound; in practice, **DuoGPT**’s savings on SRAM accesses can be even higher.

Model size results reported in Table 3(b) and Table 4 include the full parameter count, including all model parameters beyond the transformer blocks (e.g., embedding tables and output projection layers).

C Extra Experimental Results

C.1 Extra Performance Comparisons with Structured Pruning Baselines

In Table 5, we provide additional comparison results against structured pruning baselines on LLaMA-2-13B. All structured pruning baselines are pruned with 30% weight-only sparsity. The structured pruning methods use 256 calibration samples (except 2SSP, which uses 32 samples following its original setup) with 2048 token sequences. We employ 50% dual-sparsity for **DuoGPT**. The results demonstrate that **DuoGPT** achieves better performance across speedup, model size, and accuracy metrics. Notably, 50% dual-sparse **DuoGPT** achieves even higher speedup ($1.51\times$) compared to the 30% structured pruned models, where the fastest baseline achieves $1.41\times$ speedup.

Table 5: Extra comparison result with other structured pruning methods on LLaMA-2-13B.

Model	Method	Model Size	Speedup	Wiki2(↓)	PIQA	HellaSw	ARC-E	ARC-C	WinoG	Avg(↑)
Llama2-13B	ShortGPT	9.21B/13.02B	$1.41\times$	39.78	69.80	57.94	52.86	35.75	69.06	57.08
	2SSP	9.21B/13.02B	$1.29\times$	9.07	76.66	70.51	65.28	38.74	68.90	64.02
	BlockPruner	9.21B/13.02B	$1.38\times$	9.67	73.94	64.39	61.87	37.37	66.61	60.84
	DuoGPT	6.67B/13.02B	$1.51\times$	7.17	<u>75.63</u>	<u>69.56</u>	69.95	41.13	68.11	64.88

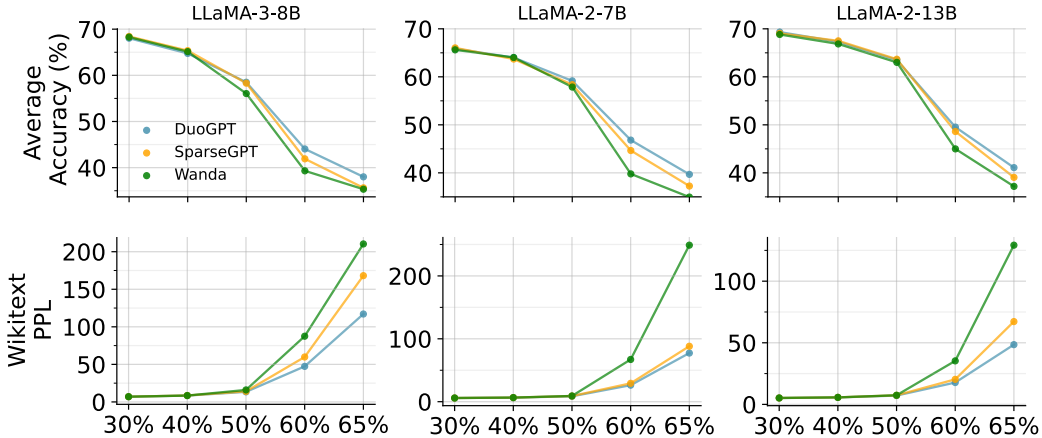


Figure 4: Mean zero-shot accuracy and perplexity on LLaMA-3-8B, LLaMA-2-7B, and LLaMA-2-13B. The results are reported across different dual-sparsity levels. The perplexity is reported for wikitext-2 dataset and the accuracy results are averaged across 7 tasks.

C.2 Performance across Different Dual-Sparsity Levels

We provide a comprehensive visualization of our approach’s performance compared to baselines across different dual-sparsity levels in Figure 4. The two baselines are dual-sparse variants of SparseGPT and Wanda, using identical setups as those in Table 1. We plot perplexity (PPL) performance on the WikiText-2 dataset and mean zero-shot accuracy across 7 tasks. We visualize sparsity levels of 30%, 40%, 50%, 60%, and 65%. Beyond 65% dual-sparsity, performance degrades too significantly from the dense model to be meaningful, while below 30%, efficiency gains are insufficient to justify the complexity.

Several key trends emerge from Figure 4. At low dual-sparsity regimes ($<50\%$), performance differences between methods are minimal regardless of calibration settings. However, the benefits of activation sparsity-aware calibration become pronounced at high dual-sparsity regimes ($\geq 50\%$). For example, Wanda’s perplexity performance (without activation sparsity-aware calibration) begins to degrade exponentially beyond 60% dual-sparsity. The SparseGPT baseline partially recovers this performance loss through activation sparsity-aware calibration, though these benefits diminish at even higher sparsity levels (65%).

DuoGPT addresses this limitation through output residual compensation, using dense model outputs to compensate for information loss during sparse calibration. This enables superior performance

in extreme dual-sparsity regions such as 65%. This trend is consistently observed across different metrics (perplexity and zero-shot accuracy) and model scales, as visualized in Figure 4. Detailed perplexity and zero-shot accuracy results for each task are provided in Table 6 for reference.

Table 6: Detailed zero-shot accuracy results for LLaMA-3-8B, LLaMA-2-7B, and LLaMA-2-13B across different dual-sparsity levels. ARC-E stands for ARC-Easy, ARC-C stands for ARC-Challenge, WinoG stands for WinoGrande, and OBQA stands for OpenBookQA. Bold and underlined values indicate the best and second-best results, respectively.

Model	Method	Wiki2(↓)	PIQA	HellaSwag	ARC-E	ARC-C	WinoG	BoolQ	OBQA	Avg(↑)
LLaMA-3-8B	SparseGPT	7.00	<u>78.94</u>	<u>77.20</u>	<u>76.43</u>	<u>49.06</u>	<u>72.85</u>	<u>81.07</u>	43.40	<u>68.42</u>
	Wanda	6.97	79.27	77.10	75.63	50.17	71.82	80.18	43.80	68.28
	DuoGPT	6.96	78.35	77.08	<u>75.84</u>	48.55	<u>72.14</u>	<u>80.92</u>	<u>43.40</u>	68.04
W _{30%} X _{30%}	SparseGPT	8.57	<u>77.26</u>	<u>73.07</u>	<u>71.68</u>	44.97	70.88	78.81	40.80	65.35
	Wanda	8.46	76.28	72.61	71.42	46.59	<u>69.22</u>	77.58	42.20	65.13
	DuoGPT	8.47	<u>76.88</u>	73.21	70.88	44.28	68.11	78.59	41.20	64.74
W _{50%} X _{50%}	SparseGPT	14.05	72.91	61.96	62.29	37.71	62.43	74.62	36.20	58.30
	Wanda	15.98	71.44	57.01	59.34	35.84	61.33	70.80	36.60	56.05
	DuoGPT	13.41	73.72	<u>61.88</u>	62.96	<u>36.77</u>	64.72	<u>74.25</u>	35.20	58.50
W _{60%} X _{60%}	SparseGPT	59.85	<u>58.32</u>	<u>33.96</u>	<u>39.39</u>	<u>23.38</u>	<u>52.25</u>	<u>58.29</u>	27.80	41.91
	Wanda	87.50	56.64	30.64	35.02	20.82	50.99	55.29	25.80	39.31
	DuoGPT	47.33	61.32	36.95	42.76	24.83	52.57	62.29	<u>27.60</u>	44.05
W _{65%} X _{65%}	SparseGPT	59.85	<u>54.35</u>	<u>28.22</u>	29.76	21.33	<u>50.99</u>	<u>38.56</u>	<u>26.20</u>	35.63
	Wanda	87.50	53.37	28.18	28.79	22.44	50.51	38.04	26.00	35.33
	DuoGPT	47.33	56.75	28.91	32.74	20.56	51.46	49.48	26.20	38.01
LLaMA-2-7B	SparseGPT	5.87	78.45	74.97	73.57	45.39	69.14	<u>77.06</u>	43.60	66.03
	Wanda	<u>5.85</u>	<u>78.29</u>	75.30	<u>73.32</u>	46.16	67.56	76.30	42.40	65.62
	DuoGPT	5.84	77.91	74.66	73.06	<u>45.65</u>	<u>68.43</u>	77.34	44.00	<u>65.86</u>
W _{30%} X _{30%}	SparseGPT	6.52	<u>77.04</u>	71.80	70.12	42.75	66.30	76.12	41.80	63.70
	Wanda	6.50	77.26	72.01	69.99	<u>42.92</u>	67.64	75.60	42.80	64.03
	DuoGPT	6.48	76.61	71.57	71.13	43.34	<u>67.17</u>	76.18	<u>42.00</u>	<u>64.00</u>
W _{40%} X _{40%}	SparseGPT	8.98	74.43	63.89	<u>62.67</u>	35.75	<u>63.46</u>	<u>71.56</u>	36.80	<u>58.37</u>
	Wanda	9.13	73.01	62.30	61.66	35.32	62.75	70.24	39.80	57.87
	DuoGPT	8.58	<u>73.83</u>	<u>63.74</u>	64.18	<u>35.67</u>	65.35	72.54	<u>38.80</u>	59.16
W _{50%} X _{50%}	SparseGPT	29.21	<u>61.15</u>	<u>38.20</u>	<u>42.38</u>	<u>24.57</u>	<u>52.41</u>	63.67	<u>30.20</u>	<u>44.65</u>
	Wanda	67.13	55.93	29.67	34.01	22.70	50.59	58.59	27.00	39.78
	DuoGPT	26.27	63.49	41.54	45.03	26.96	56.59	<u>63.30</u>	30.80	46.82
W _{60%} X _{60%}	SparseGPT	88.17	53.43	29.04	29.25	22.53	49.49	51.38	25.60	37.25
	Wanda	248.9	52.18	28.10	28.16	24.32	50.20	38.26	23.40	34.95
	DuoGPT	77.34	55.44	29.80	32.62	21.59	50.83	58.90	28.60	39.68
W _{65%} X _{65%}	SparseGPT	5.20	79.71	78.62	75.88	49.66	71.90	80.83	46.80	69.06
	Wanda	5.22	79.43	78.56	75.80	49.49	71.35	80.12	47.00	68.82
	DuoGPT	5.18	80.36	78.57	75.46	49.15	<u>71.74</u>	81.77	48.20	69.32
LLaMA-2-13B	SparseGPT	5.70	79.33	<u>76.45</u>	73.36	47.10	71.98	81.10	43.20	67.50
	Wanda	5.69	78.56	76.73	72.43	46.42	68.51	79.88	45.40	66.85
	DuoGPT	5.66	<u>78.67</u>	75.97	<u>73.02</u>	<u>47.01</u>	<u>70.80</u>	<u>80.83</u>	<u>44.80</u>	<u>67.30</u>
W _{30%} X _{30%}	SparseGPT	7.39	<u>76.28</u>	69.14	<u>68.52</u>	41.81	68.67	79.30	42.00	63.67
	Wanda	7.41	77.20	69.15	67.93	40.19	66.54	77.28	42.80	63.01
	DuoGPT	7.17	<u>76.12</u>	69.43	69.49	<u>40.96</u>	<u>67.40</u>	<u>77.83</u>	43.40	<u>63.52</u>
W _{40%} X _{40%}	SparseGPT	20.33	64.74	43.20	49.83	28.33	<u>56.27</u>	68.13	29.80	48.61
	Wanda	35.33	63.00	36.90	44.65	25.09	52.17	62.72	30.40	44.99
	DuoGPT	17.74	66.38	45.74	50.29	28.41	57.06	<u>67.25</u>	31.60	49.53
W _{50%} X _{50%}	SparseGPT	67.21	54.90	29.76	30.77	20.90	49.41	61.87	25.80	39.06
	Wanda	129.22	52.50	29.31	29.04	22.44	51.07	50.70	25.20	37.18
	DuoGPT	48.49	58.16	31.76	34.34	<u>22.35</u>	51.62	62.20	27.20	41.09

C.3 Sensitivity Analysis on the Calibration Set Size and Sequence Length

We present an ablation study to analyze the role of the C4 calibration dataset. We focus on DuoGPT with 50% dual-sparsity using the LLaMA-2-7B model. For the calibration set size study, we fix the sequence length to 2048 tokens. For the calibration sequence length study, we fix the number of calibration samples to 128.




Figure 5: The effect of the C4 calibration set size and sequence length on PPL of WikiText2 dataset for LLaMA-2-7B.

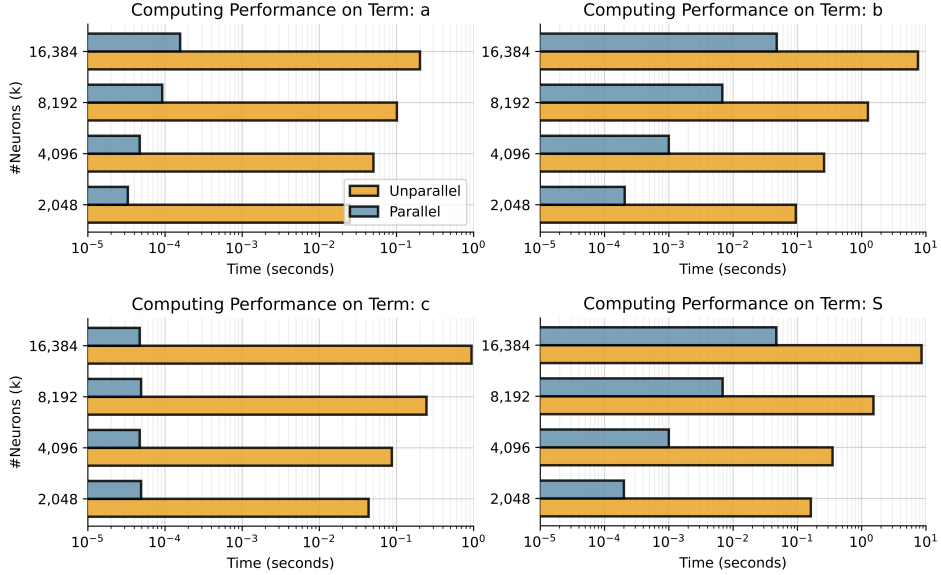


Figure 6: Latency visualization of our algorithm under different size of input channel k .

Figure 5 (left) shows the effect of varying the number of calibration samples on WikiText2 perplexity. The results demonstrate that at least 128 calibration samples provide reasonable performance for our calibration procedure.

We next explore the effect of different sequence lengths in the calibration dataset. Interestingly, we find that beyond a sequence length of 2048 tokens, the perplexity does not continue to decrease, as shown in Figure 5 (right). We conclude that for activation sparsity-aware calibration, a sequence length of 2048 tokens is sufficient to achieve good perplexity performance.

C.4 Algorithm Efficiency

Finally, we provide a visualization of the GPU runtime for our efficient **DuoGPT** implementation (Equation 7) compared to the unparallelized implementation (Equation 6). We measure latency on a single A100 80GB GPU using PyTorch 2.4.0. We provide 10 warm up runs. The token dimension (m) is fixed as 2048. Figure 6 compares the latency for computing a, b, c, and the overall pruning score S . Owing to highly optimized CUDA kernels, our vectorized precomputation of all quantities completes in approximately 1ms for a 4096×4096 layer with $m = 2048$ tokens, achieving roughly $350\times$ speedup over the unparallelized implementation. Note that we do not compare against the naive implementation that follows the optimal pruning order, as it is prohibitively slow and would not yield meaningful comparisons.